

ChurnAnalysis-ANN

January 2, 2023

0.0.1 Problem statement

Pinnacle Development Bank (PDB) is a leading retail bank in the country. The bank is facing several challenges in recent time such as Customer Churn, increment in NPA (Non-performing assets), low customer acquisition and high grievances, etc.

The cost of acquiring new customers is five times higher than the cost of retaining existing customers. Thus it is imperative to prevent customer churn. Mr. Anupam who is the head of sales decided to take help from the data science team to build a Machine learning model which will be capable to predict customer churn in advance.

****To prevent customer churn it is necessary to know which customer is going to churn. Unfortunately, domain experts can help to a very little extent in this case. We do not have the capability to predict anything in the future. With the development of computer science and statistical approaches, we can find the pattern, which would be helpful to predict it in advance.***

0.0.2 Feature Details

CustomerID - Customer identification Unique Id for each customer.

DateOfBirth - Date of birth of the customer

Gender - Gender of the customer

City - City where customer lives in

AccountBalance - Current account balance available in the account

HavingFD - If customer has FD, 0- No DF, 1- Yes FD(s) as there.

HavingCC - If customer has Credit Card, 0- No DF, 1- Yes FD(s) as there.

CIBIL_Score - Lies between 300-900 signifies customer's creditworthiness.

HavingLoan - If customer has Loan, 0- No DF, 1- Yes Loan as there.

RV - Total quantitative relationship values with the bank. Including all (AccountBalance, FD, Credit Card, Loan)

BranchType - Type of branch-Rural, Semi-urban and Urban based on the geographihcal and Popula-tion)

Churn - value is 0 and 1. 1 means customer has left the bank 0 mean did not churn)

```
[2]: #importing required libraries-
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[3]: #Setting the float format(don't want in scientific notation)
pd.set_option('display.float_format', lambda x: '%.3f' % x)
```

```
[4]: #Visualization format
sns.set_theme(context='notebook', style='darkgrid', palette='deep',
font='sans-serif',
font_scale=1, color_codes=True, rc=None)
```

```
[5]: #Don't want to print warning error messages
import warnings
warnings.filterwarnings('ignore')
```

```
[6]: #Importing the data
df = pd.read_csv('/kaggle/input/churn/churn.csv')
```

```
[7]: #Checking first 5 records
df.head()
```

```
[7]:
```

	CustomerID	DateOfBirth	Gender	City	AccountBalance	HavingFD	\
0	C7975142	01-02-1930	M	SALODARIYA	469	0	
1	C7496772	15-07-1930	F	JHAJJAR	28544	1	
2	C7541191	21-11-1930	M	KALYAN	107470	1	
3	C6531683	09-04-1932	M	MARGAO	25829	0	
4	C4715364	05-05-1932	M	CURCHOREM	216	0	

	HavingCC	CIBIL_Score	HavingLoan	RV	BranchType	Churn
0	0	833	0	163076	Rural	0
1	0	798	0	131878	Rural	0
2	0	602	1	161612	Rural	1
3	0	594	1	57874	Rural	0
4	0	808	1	184943	Rural	0

```
[8]: #Shape
df.shape
```

```
[8]: (332008, 12)
```

```
[9]: #Checking basic info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 332008 entries, 0 to 332007
```

Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	CustomerID	332008 non-null	object
1	DateOfBirth	331955 non-null	object
2	Gender	332008 non-null	object
3	City	332008 non-null	object
4	AccountBalance	332008 non-null	int64
5	HavingFD	332008 non-null	int64
6	HavingCC	332008 non-null	int64
7	CIBIL_Score	332008 non-null	int64
8	HavingLoan	332008 non-null	int64
9	RV	332008 non-null	int64
10	BranchType	331962 non-null	object
11	Churn	332008 non-null	int64

dtypes: int64(7), object(5)

memory usage: 30.4+ MB

```
[10]: #Checking data-types
df.dtypes
```

```
[10]: CustomerID      object
DateOfBirth      object
Gender           object
City            object
AccountBalance    int64
HavingFD         int64
HavingCC         int64
CIBIL_Score      int64
HavingLoan       int64
RV              int64
BranchType       object
Churn            int64
dtype: object
```

We have 5 Object(Non Numerical) and 7 Numerical Datatype

```
[11]: #Descibring the data
df.describe()
```

```
[11]:
```

	AccountBalance	HavingFD	HavingCC	CIBIL_Score	HavingLoan	\
count	332008.000	332008.000	332008.000	332008.000	332008.000	
mean	103598.116	0.364	0.309	599.667	0.448	
std	685848.257	0.481	0.462	173.392	0.497	
min	0.000	0.000	0.000	300.000	0.000	
25%	4617.750	0.000	0.000	449.000	0.000	
50%	16188.000	0.000	0.000	600.000	0.000	
75%	54066.000	1.000	1.000	750.000	1.000	

max	164489264.000	1.000	1.000	899.000	1.000
-----	---------------	-------	-------	---------	-------

	RV	Churn
count	332008.000	332008.000
mean	198624.251	0.219
std	688073.388	0.413
min	0.000	0.000
25%	69274.500	0.000
50%	133492.500	0.000
75%	194690.250	0.000
max	164608203.000	1.000

0.0.3 Observations:-

- We have 25% of customers whose account balance is more than 54066.00
- Average CIBIL Score is 600 which shows we have a good set of customers who is financially aware and pay all the due EMI/Credit card on time.
- 25% of the customers have 750 or above CIBIL Score, It is very important to signify one's financial health.
- The standard deviation of RV (relationship value) is 688073.00 which is very high and it shows that few customers are very rich and others are poor. There is a significant difference in wealth distribution.

```
[12]: #Checking null/missing values
df.isnull().sum()
```

```
[12]: CustomerID      0
DateOfBirth      53
Gender            0
City              0
AccountBalance    0
HavingFD          0
HavingCC          0
CIBIL_Score       0
HavingLoan        0
RV                0
BranchType       46
Churn             0
dtype: int64
```

There are some missing values in *DateOfBirth* and *BranchType* column which I'll impute during EDA

```
[13]: df.dtypes
```

```
[13]: CustomerID      object
      DateOfBirth    object
      Gender          object
      City            object
      AccountBalance  int64
      HavingFD        int64
      HavingCC        int64
      CIBIL_Score     int64
      HavingLoan      int64
      RV              int64
      BranchType      object
      Churn            int64
      dtype: object
```

0.0.4 Feature Engineering & Exploratory Data Analysis(EDA)

Feature Engineering is a process to use domain knowledge in order to derive a new feature or modify existing feature to extract maximum information from it.

In Feature Engineering I will perform below operations:-

1. *I will fetch Age using DateofBirth Column.*
2. *I will categories customers based on their Age.*

Exploratory data analysis is an approach of analyzing data sets to summarize their main characteristics, often using statistical graphics and other data visualization methods

I will perform below steps in EDA Section:-

1. *Duplicate data finding*
2. *Analyzing missing values(Missing value imputation).*
3. *Exploring all the features(Categorical & Numerical separately).*
4. *Finding Outliers.*
5. *Finding relationship(Correlation)*

0.0.5 A. Feature Engineering

```
[14]: df.columns
```

```
[14]: Index(['CustomerID', 'DateOfBirth', 'Gender', 'City', 'AccountBalance',
          'HavingFD', 'HavingCC', 'CIBIL_Score', 'HavingLoan', 'RV', 'BranchType',
          'Churn'],
          dtype='object')
```

We can easily derive Age from the given DateOfBirth

```
[15]: #We need datetime package for that
      from datetime import datetime
```

```
today = datetime.now()
```

```
[16]: df['DateOfBirth'] = pd.to_datetime(df['DateOfBirth']) #Convering object to Datetime
      df['Age'] = (today - df['DateOfBirth']).apply(lambda x: np.round(x.days / 365,0)) #Getting age column
```

Now based on the age I will categorize as *minor < 18, Young , 18- 35 Young-Adult, 36-60 Adult, 60>Senior Citizen*

```
[17]: age_categories = ['Minor', 'Young-Adult', 'Adult', 'Senior Citizen']
      age_bins = [0, 18, 35, 60, 100]
```

```
[18]: df['Category'] = pd.cut(df['Age'], age_bins, labels=age_categories)
```

```
[19]: df.head()
```

```
[19]: CustomerID DateOfBirth Gender      City AccountBalance HavingFD \
0    C7975142  1930-01-02      M  SALODARIYA           469         0
1    C7496772  1930-07-15      F    JHAJJAR          28544         1
2    C7541191  1930-11-21      M    KALYAN          107470         1
3    C6531683  1932-09-04      M    MARGAO           25829         0
4    C4715364  1932-05-05      M   CURCHOREM            216         0

      HavingCC  CIBIL_Score  HavingLoan      RV BranchType  Churn   Age \
0           0           833           0  163076      Rural    0  93.000
1           0           798           0  131878      Rural    0  93.000
2           0           602           1  161612      Rural    1  92.000
3           0           594           1   57874      Rural    0  90.000
4           0           808           1  184943      Rural    0  91.000

      Category
0  Senior Citizen
1  Senior Citizen
2  Senior Citizen
3  Senior Citizen
4  Senior Citizen
```

Now since we get the age, I will drop *DateOfBirth*

```
[20]: #before that I will create a copy to be in safer side.
      df_new = df.copy()
```

```
[21]: df_new.drop( 'DateOfBirth', axis = 1, inplace = True)
```

```
[22]: df_new.head()
```

```
[22]: CustomerID Gender      City AccountBalance HavingFD HavingCC \
0    C7975142      M  SALODARIYA           469         0         0
1    C7496772      F    JHAJJAR          28544         1         0
2    C7541191      M    KALYAN          107470         1         0
3    C6531683      M    MARGAO           25829         0         0
4    C4715364      M  CURCHOREM           216         0         0

      CIBIL_Score HavingLoan      RV BranchType Churn  Age      Category
0           833           0  163076      Rural    0 93.000  Senior Citizen
1           798           0  131878      Rural    0 93.000  Senior Citizen
2           602           1  161612      Rural    1 92.000  Senior Citizen
3           594           1   57874      Rural    0 90.000  Senior Citizen
4           808           1  184943      Rural    0 91.000  Senior Citizen
```

0.0.6 B. Exploratory Data Analysis(EDA)

1. Finding duplicate data

```
[23]: df_new.duplicated().sum()
```

```
[23]: 2422
```

```
[24]: df_new.duplicated(subset= 'CustomerID').sum()
```

```
[24]: 47301
```

I will drop these duplicate customers.

```
[25]: df_new.drop_duplicates(subset= 'CustomerID', inplace= True)
```

I will also drop CustomerID and City because this is a static feature and doesn't have any significance in our analysis

```
[26]: df_new.drop( ['CustomerID','City'], axis = 1, inplace = True)
```

2. Analyzing missing values

```
[27]: df_new.isnull().sum()
```

```
[27]: Gender           0
AccountBalance      0
HavingFD            0
HavingCC            0
CIBIL_Score         0
HavingLoan          0
RV                  0
BranchType          33
Churn               0
Age                 48
Category            48
```

dtype: int64

For Age I will impute mean value(as it is numerical feature) and for BranchType and use mode imputation

```
[28]: df_new['Age'] = df_new['Age'].fillna(value= df_new['Age'].mean())
      df_new['BranchType'] = df_new['BranchType'].fillna(value= df_new['BranchType'].
      ↪mode()[0])
```

```
[29]: #Defining the category after missing value imputation
      df_new['Category'] = pd.cut(df_new['Age'], age_bins, labels=age_categories)
```

```
[30]: df_new.isnull().sum()
```

```
[30]: Gender          0
      AccountBalance  0
      HavingFD        0
      HavingCC        0
      CIBIL_Score     0
      HavingLoan      0
      RV              0
      BranchType      0
      Churn           0
      Age             0
      Category        0
      dtype: int64
```


3. Feature Exploration.

```
[97]: df_new2 = df_new.copy()
      df_new2['Churned'] = np.where(df_new2['Churn'] == 1, 'Yes', 'No')
```

```
[98]: df_new2.drop('Churn', axis = 1, inplace = True)
      df_new2.rename(columns = {'Churned' : 'Churn'}, inplace = True)
```

```
[99]: df_new2['Churn'].value_counts().plot(kind = 'bar')
      df_new2['Churn'].value_counts(normalize= True)
```

```
[99]: No      0.781
      Yes    0.219
      Name: Churn, dtype: float64
```


It is a highly imbalanced dataset, 78% is not churned and 21% churned


Categorical features


```
[100]: categorical_features = [feature for feature in df_new.columns if
    ↳ df_new[feature].dtypes == 'O' or df_new[feature].dtypes == 'category']

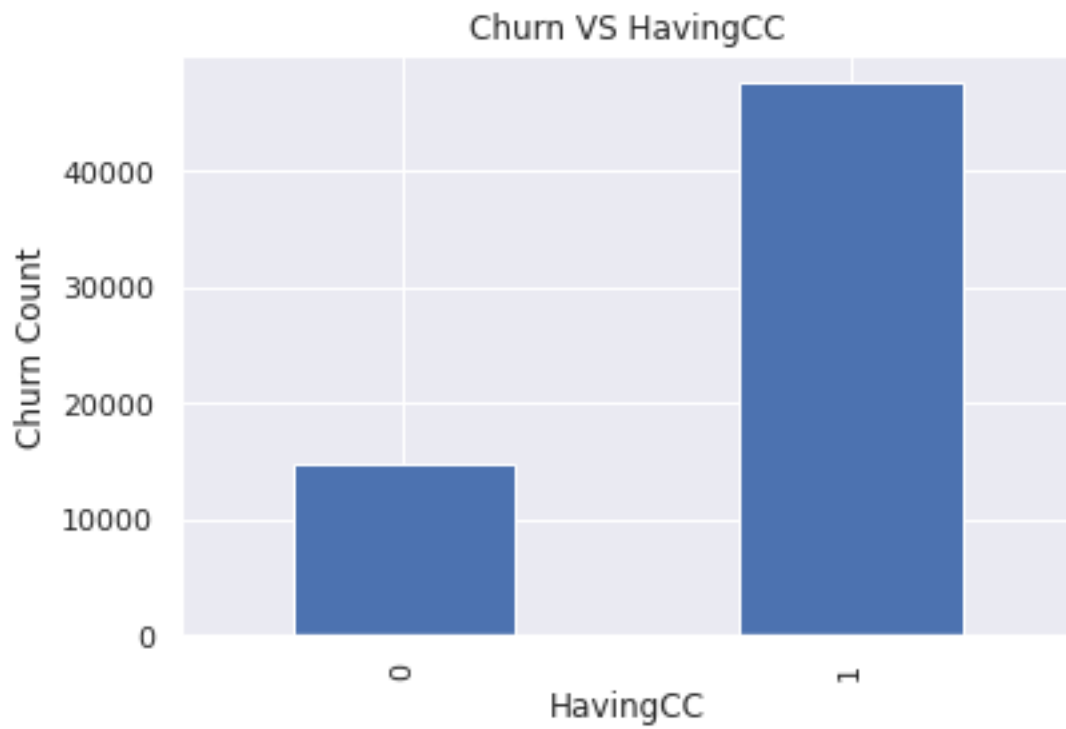

[101]: categorical_features.extend(['HavingFD', 'HavingCC', 'HavingLoan'])

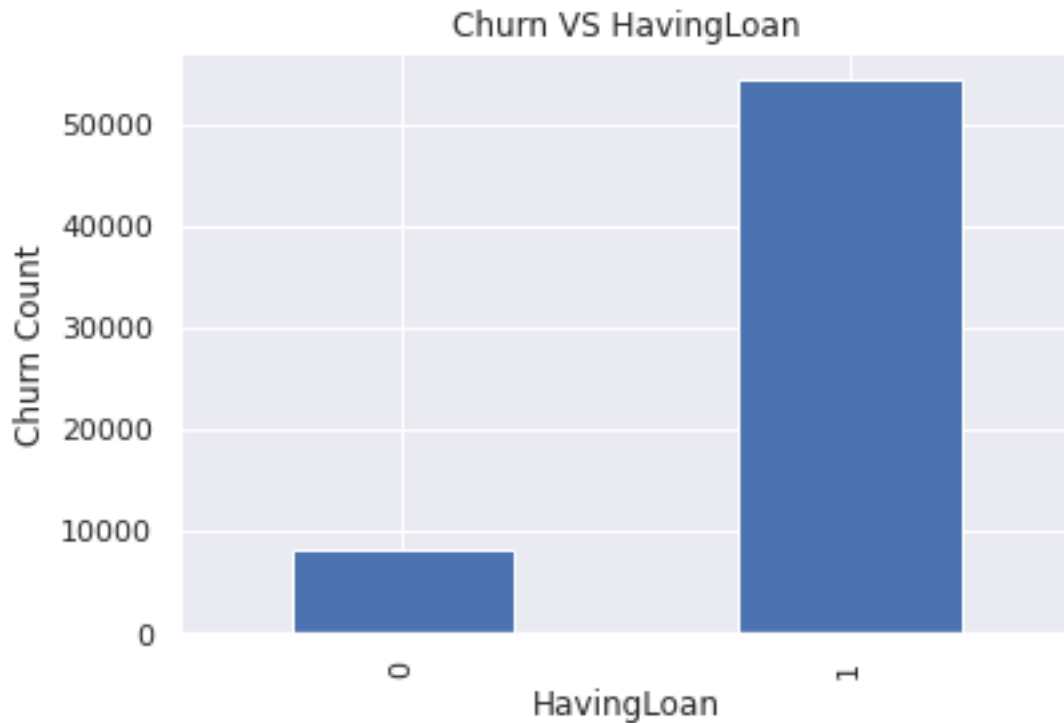
[107]: for i in categorical_features:
    df_new2[(df_new2.Churn == 'Yes')].groupby(i)['Churn'].count().plot.bar()
    plt.xlabel(i)
    plt.ylabel('Churn Count')
    plt.title('Churn VS {}'.format(i))
    plt.show()
```













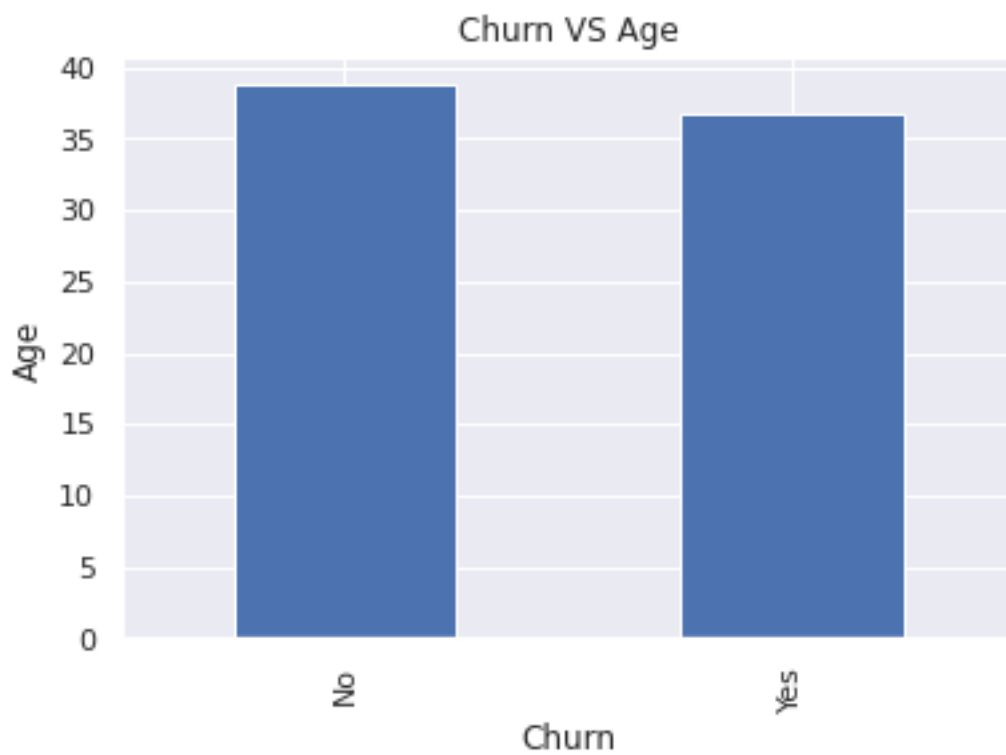

Observations:-

- *Male customer has churned more than female*
- *Customer residing in Urban Area has churned more than Rural or Semi-urban. This could be because the Urban customer may have more banks to open accounts.*
- *Young-Adult and Adult has more tendency to churn compared to Other categories. This is because they are mostly salary holders and many other banks approach them via mail/sms/calls.*
- *Customer who has FD, Credit Card or Loan has less tends to churn. Bank should try to sell them these product to prevent Churn.*

Numerical features

```
[55]: for i in df_new2:
        if i not in categorical_features and i != 'Churn':
            df_new2.groupby('Churn')[i].mean().plot.bar()
            plt.xlabel('Churn')
            plt.ylabel(i)
            plt.title('Churn VS {}'.format(i))
            plt.show()
```








Observations:-


- Customer who has churned had a very low average account balance/RV.
- There is no significant impact of CIBIL Score in Churning.
- Average age of the customer who churned is around 32-33 years(As shown in the category- Adult and Young-Adult)

4. Outliers Check.

```
[56]: #df_new2 = df_new.copy()
for i in df_new2.columns:
    if i not in categorical_features and i != 'Churn':
        sns.boxplot(df_new2[i])
        plt.xlabel(i)
        plt.show()
```







Observations:-


- *Account Balance and RV(Relationship Value) has many outliers. This is because many customers have a very high account balances.*
- *CIBIL Score has no outliers because it falls between 300-900.*
- *Age has slight outliers because of Senior Citizen Customers and few Minor Customers.*

Note:- *Despite so many outliers I will not remove these outliers because after removing it will not be a true representative of the actual data. Since we are dealing with bank data, an imbalance dataset is a common phenomenon.*

5. Correlation.

```
[57]: plt.figure(figsize=(10,6))
      sns.heatmap(df_new.corr(), annot= True,cmap="crest")
```

```
[57]: <AxesSubplot:>
```



```
[38]: #Checking absolute Correlation
df_new.corr()['Churn'].abs().sort_values(ascending = False)
```

```
[38]: Churn          1.000
      HavingCC      0.523
      HavingLoan    0.449
      Age           0.093
      HavingFD      0.079
      AccountBalance 0.048
      RV            0.034
      CIBIL_Score   0.003
      Name: Churn, dtype: float64
```

Observations:-

- Credit Card and Loan Customer has a positive Correlation. This means if someone has a Credit card or Loan has tend to churn more.
- Account balance and Age have very slight correlations.

0.0.7 Model Building

Now turn to building an efficient and accurate(as much as possible) model which will help to predict if a customer is going to churn or not.

There are several steps during model building:-

1. Creating Dummy variable(for categorical features) & Definig X and y variable.
2. Train and Test splitting.
3. Selecting a proper model(Out of few model tested).
4. Getting proper model parameters using Cross Validation Technique.
5. Measuring the Accuracy.
6. Deploying the model into production

1. Creating Dummy variable & Definig X and y variable

```
[39]: df2 = pd.get_dummies(df_new, columns= categorical_features, drop_first= True)
```

```
[40]: X = df2.drop('Churn', axis = 1)
      y = df2['Churn']
```

2. Train and Test splitting

Since we have highly imbalance data, I will use stratified splitting to get actual ratio of Churn.

```
[41]: from sklearn.model_selection import train_test_split

      x_train,x_test,y_train,y_test = train_test_split(X,y, test_size= 0.3,
      ↪random_state= 100, stratify= y)
```

3. Train and Test splitting

I will build ANN Model using Tensorflow

```
[42]: #ANN architecture
      import tensorflow
      from tensorflow.keras import Sequential
      from tensorflow.keras.layers import Dense
```

```
[60]: model = Sequential()
      model.add(Dense(11,activation='relu',input_dim=13))
      model.add(Dense(7,activation='relu'))
      model.add(Dense(1,activation='sigmoid')) ##For Binary activation function =
      ↪sigmoid
```

```
[61]: model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 11)	154
dense_7 (Dense)	(None, 7)	84

```

-----
dense_8 (Dense)                (None, 1)                8
=====
Total params: 246
Trainable params: 246
Non-trainable params: 0
-----

```

```
[62]: model.compile(optimizer='Adam',loss='binary_crossentropy',metrics=['accuracy'])
```

```
[82]: history = model.
      ↪fit(x_train,y_train,batch_size=50,epochs=50,verbose=1,validation_split=0.2)
```

```

Epoch 1/50
3189/3189 [=====] - 6s 2ms/step - loss: 14.7528 -
accuracy: 0.7128 - val_loss: 29.0681 - val_accuracy: 0.7831
Epoch 2/50
3189/3189 [=====] - 6s 2ms/step - loss: 11.2415 -
accuracy: 0.7150 - val_loss: 22.5441 - val_accuracy: 0.7830
Epoch 3/50
3189/3189 [=====] - 6s 2ms/step - loss: 8.3382 -
accuracy: 0.7119 - val_loss: 5.5955 - val_accuracy: 0.5365
Epoch 4/50
3189/3189 [=====] - 7s 2ms/step - loss: 4.1849 -
accuracy: 0.7099 - val_loss: 2.2093 - val_accuracy: 0.7626
Epoch 5/50
3189/3189 [=====] - 7s 2ms/step - loss: 3.4722 -
accuracy: 0.7122 - val_loss: 0.7640 - val_accuracy: 0.7849
Epoch 6/50
3189/3189 [=====] - 7s 2ms/step - loss: 1.2846 -
accuracy: 0.7216 - val_loss: 0.8388 - val_accuracy: 0.7829
Epoch 7/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.6844 -
accuracy: 0.7441 - val_loss: 0.4376 - val_accuracy: 0.7907
Epoch 8/50
3189/3189 [=====] - 6s 2ms/step - loss: 0.5609 -
accuracy: 0.7739 - val_loss: 0.4933 - val_accuracy: 0.7828
Epoch 9/50
3189/3189 [=====] - 6s 2ms/step - loss: 0.4963 -
accuracy: 0.7804 - val_loss: 0.4924 - val_accuracy: 0.7828
Epoch 10/50
3189/3189 [=====] - 6s 2ms/step - loss: 0.4960 -
accuracy: 0.7804 - val_loss: 0.4915 - val_accuracy: 0.7828
Epoch 11/50
3189/3189 [=====] - 6s 2ms/step - loss: 0.4947 -
accuracy: 0.7804 - val_loss: 0.4901 - val_accuracy: 0.7827
Epoch 12/50
3189/3189 [=====] - 6s 2ms/step - loss: 0.5309 -

```

accuracy: 0.7798 - val_loss: 0.4939 - val_accuracy: 0.7828
Epoch 13/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.5063 -
accuracy: 0.7798 - val_loss: 0.4912 - val_accuracy: 0.7828
Epoch 14/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.5133 -
accuracy: 0.7801 - val_loss: 0.5118 - val_accuracy: 0.7828
Epoch 15/50
3189/3189 [=====] - 8s 2ms/step - loss: 0.5147 -
accuracy: 0.7804 - val_loss: 0.5114 - val_accuracy: 0.7828
Epoch 16/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.5021 -
accuracy: 0.7804 - val_loss: 0.4911 - val_accuracy: 0.7828
Epoch 17/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.4938 -
accuracy: 0.7804 - val_loss: 0.4904 - val_accuracy: 0.7828
Epoch 18/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.4946 -
accuracy: 0.7804 - val_loss: 0.4912 - val_accuracy: 0.7828
Epoch 19/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.4944 -
accuracy: 0.7804 - val_loss: 0.4903 - val_accuracy: 0.7828
Epoch 20/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.4939 -
accuracy: 0.7804 - val_loss: 0.4899 - val_accuracy: 0.7828
Epoch 21/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.4958 -
accuracy: 0.7803 - val_loss: 0.4908 - val_accuracy: 0.7828
Epoch 22/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.4935 -
accuracy: 0.7804 - val_loss: 0.4937 - val_accuracy: 0.7828
Epoch 23/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.4988 -
accuracy: 0.7804 - val_loss: 0.4941 - val_accuracy: 0.7828
Epoch 24/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.4933 -
accuracy: 0.7804 - val_loss: 0.4907 - val_accuracy: 0.7828
Epoch 25/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.5017 -
accuracy: 0.7804 - val_loss: 0.4933 - val_accuracy: 0.7828
Epoch 26/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.4938 -
accuracy: 0.7804 - val_loss: 0.4897 - val_accuracy: 0.7828
Epoch 27/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.4975 -
accuracy: 0.7803 - val_loss: 0.4901 - val_accuracy: 0.7828
Epoch 28/50
3189/3189 [=====] - 7s 2ms/step - loss: 0.4930 -

accuracy: 0.7804 - val_loss: 0.4902 - val_accuracy: 0.7828
 Epoch 29/50
 3189/3189 [=====] - 8s 2ms/step - loss: 0.4931 -
 accuracy: 0.7804 - val_loss: 0.4901 - val_accuracy: 0.7828
 Epoch 30/50
 3189/3189 [=====] - 7s 2ms/step - loss: 0.4941 -
 accuracy: 0.7804 - val_loss: 0.4908 - val_accuracy: 0.7828
 Epoch 31/50
 3189/3189 [=====] - 7s 2ms/step - loss: 0.5031 -
 accuracy: 0.7802 - val_loss: 0.4887 - val_accuracy: 0.7828
 Epoch 32/50
 3189/3189 [=====] - 7s 2ms/step - loss: 0.5029 -
 accuracy: 0.7803 - val_loss: 0.4904 - val_accuracy: 0.7828
 Epoch 33/50
 3189/3189 [=====] - 7s 2ms/step - loss: 0.4935 -
 accuracy: 0.7804 - val_loss: 0.4900 - val_accuracy: 0.7828
 Epoch 34/50
 3189/3189 [=====] - 8s 2ms/step - loss: 0.4932 -
 accuracy: 0.7804 - val_loss: 0.4923 - val_accuracy: 0.7828
 Epoch 35/50
 3189/3189 [=====] - 8s 2ms/step - loss: 0.5011 -
 accuracy: 0.7803 - val_loss: 0.4907 - val_accuracy: 0.7828
 Epoch 36/50
 3189/3189 [=====] - 8s 2ms/step - loss: 0.4939 -
 accuracy: 0.7804 - val_loss: 0.4900 - val_accuracy: 0.7828
 Epoch 37/50
 3189/3189 [=====] - 7s 2ms/step - loss: 0.4957 -
 accuracy: 0.7804 - val_loss: 0.4966 - val_accuracy: 0.7826
 Epoch 38/50
 3189/3189 [=====] - 8s 2ms/step - loss: 0.4958 -
 accuracy: 0.7804 - val_loss: 0.4902 - val_accuracy: 0.7828
 Epoch 39/50
 3189/3189 [=====] - 8s 3ms/step - loss: 0.4934 -
 accuracy: 0.7804 - val_loss: 0.4902 - val_accuracy: 0.7828
 Epoch 40/50
 3189/3189 [=====] - 9s 3ms/step - loss: 0.4938 -
 accuracy: 0.7804 - val_loss: 0.4887 - val_accuracy: 0.7828
 Epoch 41/50
 3189/3189 [=====] - 9s 3ms/step - loss: 0.5019 -
 accuracy: 0.7802 - val_loss: 0.4895 - val_accuracy: 0.7828
 Epoch 42/50
 3189/3189 [=====] - 9s 3ms/step - loss: 0.4921 -
 accuracy: 0.7804 - val_loss: 0.4890 - val_accuracy: 0.7828
 Epoch 43/50
 3189/3189 [=====] - 9s 3ms/step - loss: 0.4938 -
 accuracy: 0.7804 - val_loss: 0.4885 - val_accuracy: 0.7828
 Epoch 44/50
 3189/3189 [=====] - 9s 3ms/step - loss: 0.4906 -


```

accuracy: 0.7804 - val_loss: 0.4874 - val_accuracy: 0.7828
Epoch 45/50
3189/3189 [=====] - 9s 3ms/step - loss: 0.4916 -
accuracy: 0.7804 - val_loss: 0.4888 - val_accuracy: 0.7828
Epoch 46/50
3189/3189 [=====] - 9s 3ms/step - loss: 0.4920 -
accuracy: 0.7804 - val_loss: 0.4868 - val_accuracy: 0.7828
Epoch 47/50
3189/3189 [=====] - 10s 3ms/step - loss: 0.4916 -
accuracy: 0.7804 - val_loss: 0.4881 - val_accuracy: 0.7828
Epoch 48/50
3189/3189 [=====] - 10s 3ms/step - loss: 0.4899 -
accuracy: 0.7804 - val_loss: 0.4918 - val_accuracy: 0.7828
Epoch 49/50
3189/3189 [=====] - 10s 3ms/step - loss: 0.4901 -
accuracy: 0.7804 - val_loss: 0.4876 - val_accuracy: 0.7828
Epoch 50/50
3189/3189 [=====] - 10s 3ms/step - loss: 0.5070 -
accuracy: 0.7804 - val_loss: 0.5063 - val_accuracy: 0.7828

```

```
[83]: y_pred = model.predict(x_test)
```

```
[84]: y_pred_f = np.where(y_pred < 0.5 , 0,1)
```

5. Accuracy check with F1 score.

```
[85]: from sklearn.metrics import f1_score,accuracy_score
```

```
[86]: f1_score(y_test,y_pred_f)
```


```
[86]: 0.0
```

```
[87]: accuracy_score(y_test,y_pred_f)
```

```
[87]: 0.7808998630185101
```

```
[88]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
```

```
[88]: [<matplotlib.lines.Line2D at 0x7ffa2ffdf7d0>]
```




Plotting AUC-ROC curve:-

```
[89]: from sklearn.metrics import roc_curve

[90]: y_score = model.predict(x_test)

[91]: fpr, tpr, thresholds = roc_curve(y_test, y_score)

[92]: # Plot the ROC curve
plt.plot(fpr, tpr, 'b-', label='ANN')
plt.plot([0, 1], [0, 1], 'k--', label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='best')
plt.show()
```



Now, this model can be deployed in any local machine or cloud server as a full production grade application This can be integrated into any web application that will help bank staff to check if a customer is going to churn and take preventive measures to stop it..

0.0.8 Copyright © Gopi Nath Pandit